

OpenML in Python

OpenML is an online collaboration platform for machine learning:

- Share/reuse machine learning datasets, algorithms, models, experiments
- Well documented/annotated datasets, uniform access
- APIs in Java, R, Python*,... to download/upload everything
- Better reproducibility of experiments, reuse of machine learning models
- Works well with machine learning libraries such as scikit-learn
- Large scale benchmarking, compare to state of the art

<IPython.core.display.HTML object>

Authentication

- Create an OpenML account (free) on <http://www.openml.org>.
- After logging in, open your account page (avatar on the top right)
- Open 'Account Settings', then 'API authentication' to find your API key.

There are two ways to authenticate:

- Create a plain text file `~/.openml/config` with the line `'apikey=MYKEY'`, replacing MYKEY with your API key.
- Run the code below, replacing 'MYKEY' with your API key.

```
[67]: # If you don't keep your API key in the config file, you can also specify
      # oml.config.apikey = MYKEY
```

Data sets

We can list, select, and download all OpenML datasets

List datasets

```
[68]: datalist = oml.datasets.list_datasets() # Returns a dict
      datalist = pd.DataFrame.from_dict(datalist, orient='index') # Create a DataFrame
      print("First 10 of %s active datasets..." % len(datalist))
      datalist[:10][['did', 'name', 'NumberOfInstances',
                    'NumberOfFeatures', 'NumberOfClasses']]
```

First 10 of 2518 active datasets...

	did	name	NumberOfInstances	NumberOfFeatures	NumberOfClasses
2	2	anneal	898	39	5
3	3	kr-vs-kp	3196	37	2
4	4	labor	57	17	2
5	5	arrhythmia	452	280	13
6	6	letter	20000	17	26
7	7	audiology	226	70	24
8	8	liver-disorders	345	7	16
9	9	autos	205	26	6

10	10	lymph	148	19	4
11	11	balance-scale	625	5	3

There are many properties that we can query

```
[69]: list(datalist)
      datalist = datalist[['did', 'name', 'NumberOfInstances',
                          'NumberOfFeatures', 'NumberOfClasses']]

['name',
 'NumberOfFeatures',
 'format',
 'NumberOfNumericFeatures',
 'NumberOfClasses',
 'NumberOfSymbolicFeatures',
 'NumberOfInstances',
 'MinorityClassSize',
 'NumberOfInstancesWithMissingValues',
 'MajorityClassSize',
 'did',
 'status',
 'NumberOfMissingValues',
 'MaxNominalAttDistinctValues']
```

and we can filter or sort on all of them

```
[70]: datalist[datalist.NumberOfInstances>10000
          ].sort(['NumberOfInstances'])[:20]
```

	did	name	NumberOfInstances	\
23515	23515	sulfur	10081	
981	981	kdd_internet_usage	10108	
372	372	internet_usage	10108	
1536	1536	volcanoes-b6	10130	
4562	4562	InternetUsage	10168	
1531	1531	volcanoes-b1	10176	
1534	1534	volcanoes-b4	10190	
1459	1459	artificial-characters	10218	
1478	1478	har	10299	
1533	1533	volcanoes-b3	10386	
1532	1532	volcanoes-b2	10668	
1053	1053	jm1	10885	
1414	1414	Kaggle_bike_sharing_demand_challenge	10886	
1044	1044	eye_movements	10936	
1019	1019	pendigits	10992	
32	32	pendigits	10992	
4534	4534	PhishingWebsites	11055	
399	399	ohscal.wc	11162	
310	310	mammography	11183	
1568	1568	nursery	12958	

	NumberOfFeatures	NumberOfClasses
23515	7	-1
981	69	2
372	72	46
1536	4	5
4562	72	-1
1531	4	5
1534	4	5
1459	8	10
1478	562	6
1533	4	5
1532	4	5
1053	22	2
1414	12	822
1044	28	3
1019	17	2
32	17	10
4534	31	2
399	11466	10
310	7	2
1568	9	4

or find specific ones

```
[71]: datalist.query('name == "MagicTelescope"')
```

	did	name	NumberOfInstances	NumberOfFeatures \
1120	1120	MagicTelescope	19020	12

	NumberOfClasses
1120	2

```
[72]: datalist.query('NumberOfClasses > 50')
```

	did	name	NumberOfInstances \
190	190	mbagrade	61
1092	1092	Crash	352
1414	1414	Kaggle_bike_sharing_demand_challenge	10886
1491	1491	one-hundred-plants-margin	1600
1492	1492	one-hundred-plants-shape	1600
1493	1493	one-hundred-plants-texture	1599
4546	4546	Plants	44940
4552	4552	BachChoralHarmony	5665
40601	40601	RAM_price	333
40753	40753	delays_zurich_transport	5465575
40916	40916	HappinessRank_2015	158
41021	41021	Moneyball	1232
41022	41022	Short_Track_Speed_Skating	5125

	NumberOfFeatures	NumberOfClasses
--	------------------	-----------------

190	3	57
1092	14	307
1414	12	822
1491	65	100
1492	65	100
1493	65	100
4546	16	57
4552	17	102
40601	3	219
40753	15	4082
40916	12	157
41021	15	87
41022	27	2954

Download a specific dataset. This is done based on the dataset ID (called 'did').

```
[73]: dataset = oml.datasets.get_dataset(1120)

print("This is dataset '%s', the target feature is '%s'" %
      (dataset.name, dataset.default_target_attribute))
print("URL: %s" % dataset.url)
print(dataset.description[:500])
```

```
This is dataset 'MagicTelescope', the target feature is 'class:'
URL: https://www.openml.org/data/v1/download/54003/MagicTelescope.arff
**Author**: R. K. Bock. Major Atmospheric Gamma Imaging Cherenkov Telescope project
Donated by P. Savicky, Institute of Computer Science, AS of CR, Czech Republic
**Source**: [UCI](https://archive.ics.uci.edu/ml/datasets/magic+gamma+telescope)
**Please cite**: Bock, R.K., Chilingarian, A., Gaug, M., Hakl, F., Hengstebeck,
Methods for multidimensional event classification: a case study
```

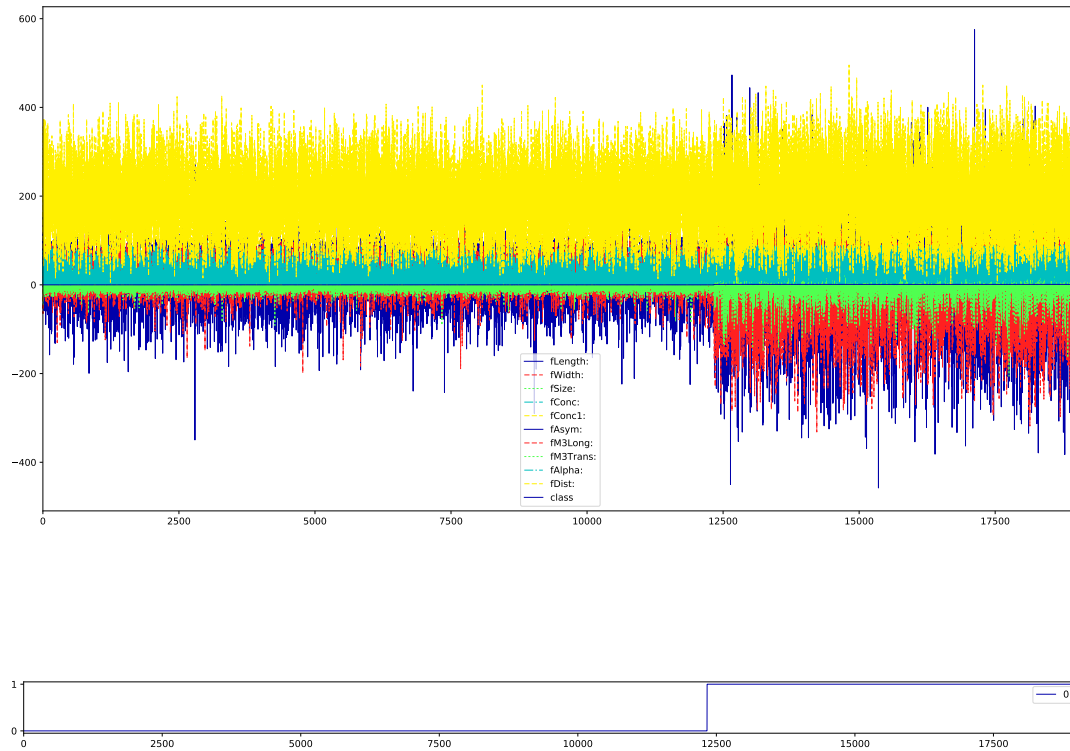
Convert the data to a DataFrame for easier processing/plotting

```
[74]: X, y, attribute_names = dataset.get_data(
      target=dataset.default_target_attribute,
      return_attribute_names=True)
magic = pd.DataFrame(X, columns=attribute_names)
magic['class'] = y
print(magic[:10])
```


	fLength:	fWidth:	fSize:	fConc:	...	fM3Trans:	fAlpha:	fDist:	class
0	28.80	16.00	2.64	0.39	...	-8.20	40.09	81.88	0
1	31.60	11.72	2.52	0.53	...	-9.96	6.36	205.26	0
2	162.05	136.03	4.06	0.04	...	-45.22	76.96	256.79	0
3	23.82	9.57	2.34	0.61	...	-7.15	10.45	116.74	0
4	75.14	30.92	3.16	0.32	...	21.84	4.65	356.46	0
5	51.62	21.15	2.91	0.24	...	9.81	3.61	238.10	0
6	48.25	17.36	3.03	0.25	...	10.59	4.79	219.09	0
7	26.79	13.76	2.55	0.42	...	-2.93	0.81	237.13	0
8	96.23	46.52	4.15	0.08	...	43.18	4.85	248.23	0
9	46.76	15.20	2.58	0.34	...	-6.68	7.88	102.25	0

```
[10 rows x 11 columns]
```

```
[75]: magic.plot(figsize=(20,10))  
      pd.DataFrame(y).plot(figsize=(20,1));
```



```
[76]: magics = magic.sample(n=1000)  
      _ = pd.scatter_matrix(magics.iloc[:100,:4], c=magics[:100]['class'], figsize=(10,10),  
                           marker='o', hist_kwds={'bins': 20},  
                           alpha=.8, cmap='plasma')
```



Train models

Train a scikit-learn model on the data manually

```
[77]: from sklearn import ensemble
      from sklearn import model_selection

      X_train, X_test, y_train, y_test = model_selection.train_test_split(X, y)
      clf = ensemble.RandomForestClassifier(n_estimators=100)
      clf.fit(X_train, y_train)
      pd.DataFrame(clf.predict(X)).plot(figsize=(20,1));
```



And evaluate

```
[78]: kfold = model_selection.StratifiedKfold(n_splits=5, shuffle=True, random_s
      results = model_selection.cross_val_score(clf, X, y, cv=kfold)
      print("Accuracy: %.3f%% (+- %.3f)" % (results.mean(), results.std()))
```

Accuracy: 0.880% (+- 0.002)

Note: You can also ask which features are categorical to do your own encoding

```
[79]: from sklearn import preprocessing
      dataset = oml.datasets.get_dataset(10) # Lymph dataset
      X, y, categorical = dataset.get_data(
          target=dataset.default_target_attribute,
          return_categorical_indicator=True)
      print("Categorical features: %s" % categorical)
      enc = preprocessing.OneHotEncoder(categorical_features=categorical)
      X = enc.fit_transform(X)
      clf.fit(X, y)
```

Categorical features: [True, True, True, True, True, True, True, True, True, False, Fa

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

Tasks

To run benchmarks consistently (also across studies and tools), OpenML offers Tasks, which include specific train-test splits and other information to define a scientific task. Tasks are typically created via the website by the dataset provider.

Listing tasks

```
[80]: task_list = oml.tasks.list_tasks(size=5000) # Get first 5000 tasks

      mytasks = pd.DataFrame(task_list).transpose()
      print("First 5 of %s tasks:" % len(mytasks))
      print(mytasks.columns)
```

First 5 of 5000 tasks:

```
Index(['MajorityClassSize', 'MaxNominalAttDistinctValues', 'MinorityClassSize',
      'NumberOfClasses', 'NumberOfFeatures', 'NumberOfInstances',
      'NumberOfInstancesWithMissingValues', 'NumberOfMissingValues',
      'NumberOfNumericFeatures', 'NumberOfSymbolicFeatures', 'cost_matrix',
      'did', 'estimation_procedure', 'evaluation_measures', 'name',
      'number_samples', 'quality_measure', 'source_data',
```

```

'source_data_labeled', 'status', 'target_feature',
'target_feature_event', 'target_feature_left', 'target_feature_right',
'target_value', 'task_type', 'tid', 'ttid'],
dtype='object')

```

```

[81]: mytasks = mytasks[['tid', 'did', 'name', 'task_type', 'estimation_procedure',
print(mytasks.head())

```

	tid	did	name	task_type	estimation_procedure	\
2	2	2	anneal	Supervised Classification	10-fold Crossvalidation	
3	3	3	kr-vs-kp	Supervised Classification	10-fold Crossvalidation	
4	4	4	labor	Supervised Classification	10-fold Crossvalidation	
5	5	5	arrhythmia	Supervised Classification	10-fold Crossvalidation	
6	6	6	letter	Supervised Classification	10-fold Crossvalidation	

	evaluation_measures
2	predictive_accuracy
3	NaN
4	predictive_accuracy
5	predictive_accuracy
6	NaN

Search for the tasks you need

```

[82]: print(mytasks.query('name=="MagicTelescope"'))

```

	tid	did	name	task_type	\
3954	3954	1120	MagicTelescope	Supervised Classification	
4659	4659	1120	MagicTelescope	Supervised Classification	
7228	7228	1120	MagicTelescope	Supervised Data Stream Classification	
10067	10067	1120	MagicTelescope	Learning Curve	

	estimation_procedure	evaluation_measures
3954	10-fold Crossvalidation	NaN
4659	10 times 10-fold Crossvalidation	predictive_accuracy
7228	Interleaved Test then Train	NaN
10067	10-fold Learning Curve	NaN

Download tasks

```

[83]: task = oml.tasks.get_task(3954)
pprint(vars(task))

```

```

{'class_labels': ['g', 'h'],
'cost_matrix': None,
'dataset_id': 1120,
'estimation_parameters': {'number_folds': '10',
                           'number_repeats': '1',
                           'percentage': '',
                           'stratified_sampling': 'true'},

```



```

'estimation_procedure': {'data_splits_url': 'https://www.openml.org/api_splits/
                        'parameters': {'number_folds': '10',
                                        'number_repeats': '1',
                                        'percentage': '',
                                        'stratified_sampling': 'true'},
                        'type': 'crossvalidation'},
'evaluation_measure': None,
'split': None,
'target_name': 'class:',
'task_id': 3954,
'task_type': 'Supervised Classification',
'task_type_id': 1}

```

Runs: Train models on tasks

We can run (many) scikit-learn algorithms on (many) OpenML tasks.

```

[84]: task = oml.tasks.get_task(3954)
      clf = ensemble.RandomForestClassifier(n_estimators=100)
      flow = oml.flows.sklearn_to_flow(clf)
      run = oml.runs.run_flow_on_task(task, flow)
      run.model

```

```

RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                        oob_score=False, random_state=38834, verbose=0,
                        warm_start=False)

```

Share the run on the OpenML server * Note: if the exact same experiment has already been done (probably by you), it will just return the run ID from the previous run

```

[85]: myrun = run.publish()
      print("Uploaded to http://www.openml.org/r/" + str(myrun.run_id))

```

Uploaded to <http://www.openml.org/r/8872325>

It also works with pipelines

```

[86]: from sklearn import pipeline, ensemble, preprocessing
      from openml import tasks, runs, datasets
      task = tasks.get_task(59)
      pipe = pipeline.Pipeline(steps=[
          ('Imputer', preprocessing.Imputer(strategy='median')),
          ('OneHotEncoder', preprocessing.OneHotEncoder(sparse=False, ha
          ('Classifier', ensemble.RandomForestClassifier())
      ])

```

```

flow = oml.flows.sklearn_to_flow(pipe)
run = oml.runs.run_flow_on_task(task, flow)
myrun = run.publish()
print("Uploaded to http://www.openml.org/r/" + str(myrun.run_id))

```

Uploaded to <http://www.openml.org/r/8872326>

All together

Train any model on any OpenML dataset and upload to OpenML in a few lines of code

```
[87]: from sklearn.linear_model import LogisticRegression
```

```

task = oml.tasks.get_task(3954)
clf = LogisticRegression()
flow = oml.flows.sklearn_to_flow(clf)
run = oml.runs.run_flow_on_task(task, flow)
run.model
myrun = run.publish()
print("Uploaded to http://www.openml.org/r/" + str(myrun.run_id))

```

```

LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=45647, solver='liblinear', tol=0.0001,
                    verbose=0, warm_start=False)

```

Uploaded to <http://www.openml.org/r/8872327>

- You can (of course) evaluate models in a loop
- E.g. for easy benchmarking:

```

[89]: import openml as oml
      from sklearn import neighbors

      for task_id in [14951, 10103, 9945]:
          task = oml.tasks.get_task(task_id)
          data = oml.datasets.get_dataset(task.dataset_id)
          clf = neighbors.KNeighborsClassifier(n_neighbors=10)
          flow = oml.flows.sklearn_to_flow(clf)
          run = oml.runs.run_flow_on_task(task, flow)
          myrun = run.publish()
          print("kNN on %s: http://www.openml.org/r/%d" % (data.name, myrun.run_id))

```

kNN on eeg-eye-state: <http://www.openml.org/r/8872328>

kNN on volcanoes-a1: <http://www.openml.org/r/8872329>

kNN on walking-activity: <http://www.openml.org/r/8872330>

Download everyone else's results on the same dataset

Check whether other people built better models on the same task by downloading their evaluations (computed on the OpenML server) and comparing directly against them.

```
[90]: myruns = oml.runs.list_runs(task=[3954],size=10000) # This can take a while
      scores = []
      for id, _ in myruns.items():
          run = oml.runs.get_run(id)
          if str.startswith(run.flow_name, 'sklearn'):
              scores.append({"flow":run.flow_name, "score":run.evaluations['pred

[94]: import seaborn as sns
      fig, ax = plt.subplots(figsize=(8, 12))
      sns.violinplot(x="score", y="flow", data=pd.DataFrame(scores), scale="width
```

